

Lecture 9

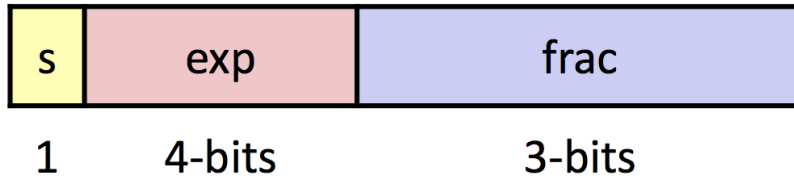
<2016-05-02 Mon>

Contents

1	Floating Point	2
1.0.1	Tiny Floating Point Example	2
1.1	Distribution of Values	2
1.2	Special Properties of the IEEE Encoding	3
1.3	Floating Point Operation	3
1.3.1	Round to Even	4
1.3.2	Floating Point Multiplication	4
1.3.3	Floating Point Addition	5
1.3.4	Floating Point in C	6
1.4	Summary	7
2	Optimization	7
2.1	Optimizing Compilers	7
2.2	Generally Useful Optimizations	8
2.2.1	Example of Memory Aliasing	8
2.3	Instruction-Level Parallelism	8
2.4	Superscalar Processor	9
2.4.1	Pipelined Functional Units	9

1 Floating Point

1.0.1 Tiny Floating Point Example



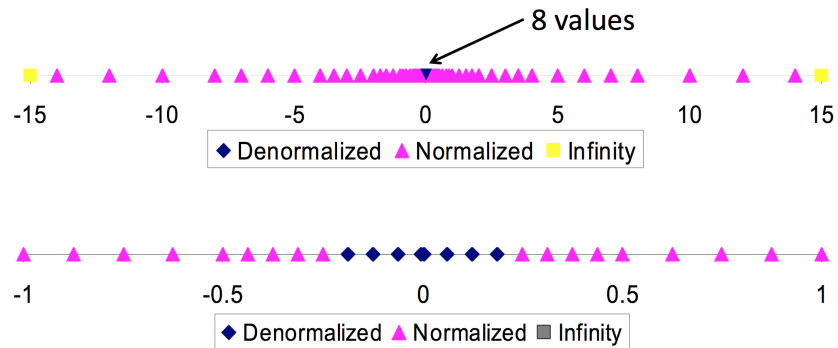
- 8-bit floating point representation
 - the sign bit is in the most significant bit
 - the next four bits are the exponent, with a bias of 7
 - the last 3 bits are the frac

state	s	exp	frac	E	value	explanation
denormalized numbers	0	0000	000	-6	0	
	0	0000	001	-6	1/512	closest to 0
	0	0000	010	-6	2/512	
	...					
	0	0000	110	-6	6/512	
	0	0000	111	-6	7/512	largest denorm
normalized numbers	0	0001	000	-6	8/512	smallest norm
	0	0001	001	-6	9/512	
	...					
	0	0110	110	-1	14/16	
	0	0110	111	-1	15/16	closest to 1 below
	0	0111	000	0	1	
	0	0111	001	0	9/8	closest to 1 above
	...					
	0	1110	110	7	224	
	0	1110	111	7	240	largest norm
	0	1111	000	n/a	inf	

1.1 Distribution of Values

- 6-bits IEEE-like format
 - e = 3 exponent bits

- $f = 2$ fraction bits
- bias is $2^{3-1} - 1 = 3$
- distribution gets denser toward 0



1.2 Special Properties of the IEEE Encoding

- floating point 0 is the same as integer 0
 - all bits 0
- can almost use unsigned integer comparison
 - must first compare sign bits
 - must consider $-0 = 0$
 - NaN (s) problematic
 - * will be greater than any other values
 - * what should comparison yield
 - otherwise OK
 - * denormalized vs normalized
 - * normalized vs infinity

1.3 Floating Point Operation

- basic idea
 - first compute exact result
 - make it fit into desired precision
 - * possibly overflow if exponent too large
 - * possibly round to fit into fraction part of float

1.3.1 Round to Even

1. Rounding Modes

mode	1.40	1.60	1.50	2.50	-1.50
towards 0	1	1	1	2	-1
round down $-\infty$	1	1	1	2	-2
round up $+\infty$	2	2	2	3	-1
nearest even	1	2	2	2	-2

- round-to-even: default rounding mode
 - hard to get any other kind without dropping into assembly
 - all other are statistically biased
 - * sum of set of positive numbers will consistently be over or under estimated
- applying to other decimal places / bit positions
 - when exactly halfway between 2 possible values
 - * **round so that least significant digit is even**

2. rounding binary numbers

- binary fractional numbers
 - "even" when least significant bit is 0
 - "half" way when bits to right of rounding position = $100 \dots_2$

Table 1: round to nearest 1/4 (2 bits right of binary point)

value	binary	rounded	action	bounded
3/32	0.00011	0.00	down	0
3/16	0.00110	0.01	up	1/4
7/8	0.11100	1.00	up	1
5/8	0.10100	0.10	down	1/2

1.3.2 Floating Point Multiplication


- $(-1)^{s_1} M 2^{E_1} \times (-1)^{s_2} M 2^{E_2}$
- exact result $(-1)^s M 2^E$
 - sign s : $s_1 \wedge s_2$

- significand M: $M1 * M2$
- exponent E: $E1 + E2$
- fixing
 - if $M \geq 2$, shift M right, increment E
 - if E out of range, overflow
 - round M to fit **frac** precision
- implementation
 - biggest chore is multiplying significand

1.3.3 Floating Point Addition

- $(-1)^{s1} M 2^{E1} + (-1)^{s2} M 2^{E2}$
- exact result: $(-1)^s M 2^2$
 - sign **s**, significand M:
 - * result of signed align & add
 - exponent E: E1
- fixing
 - if $M \geq 2$, shift M right, increment E
 - if $M < 1$, shift M left **k** positions, decrement E by **k**
 - overflow if E out of range
 - round M to fit **frac** range

Get binary points lined up



1.3.4 Floating Point in C

- C guarantees two levels
 - `float` single precision
 - `double` double precision
- conversion/casting
 - casting between `int`, `float`, and `double` changes bit representation
 - `double/float => =int`
 - * truncates fractional part
 - * like rounding towards 0
 - * not defined when out of range or NaN: generally sets to TMIN
 - `int => =double`
 - * exact conversion, as long as `int` has ≤ 53 word size
 - `int => =float`
 - * will round according to rounding mode

1. Floating Point Puzzle

```
int x;
float f;
double d;
```

condition	expression	result	explanation
f	<code>-(-f)</code>	true	
d < 0.0	<code>(d * 2 < 0.0)</code>	true	
d > f	<code>-f < -d</code>	true	
	<code>d*d >= 0.0</code>	true	
	<code>(d+f) - d == f</code>	false	add a large number d to a small number f, precision over

1.4 Summary

- IEEE Floating Point has clear mathematical properties
- represents numbers of form $M \times 2^E$
- can reason about operations independent of implementation
 - as if computed with perfect precision then rounded
- not the same as real arithmetic

2 Optimization

2.1 Optimizing Compilers

- provide efficient mapping of program to machine
 - register allocation
 - code selection and ordering (scheduling)
 - dead code elimination
 - eliminating minor inefficiencies
- don't improve asymptotic efficiency
 - up to programmers to select best overall algorithms
 - big-O savings are (often) more important than constant factors
- have difficulty overcoming "optimization blockers"
 - potential memory aliasing
 - potential procedure side-effect

2.2 Generally Useful Optimizations

optimizations that you or the compiler should do regardless of processor/compiler

2.2.1 Example of Memory Aliasing


```
void sum_row2(double *a, double *b, long n) {
    long i, j;
    for (int i = 0; i < n; ++i) {
        b[i] = 0;
        for (int j = 0; j < n; ++j)
            b[i] += a[i*n + j];
    }
}
```

```
void sum_row2(double *a, double *b, long n) {
    long i, j;
    for (int i = 0; i < n; ++i) {
        double val = 0;          /* remove memory aliasing */
        for (int j = 0; j < n; ++j)
            val += a[i*n + j];
        b[i] = val;
    }
}
```

- code updates `b[i]` on every iteration
- must consider possibilities that these updates will affect program behavior

2.3 Instruction-Level Parallelism

- need general understanding of modern processor design
 - hardware can execute multiple instructions in parallel
- performance limited by data dependencies
- simple transformations can yield dramatic performance improvement
 - compilers often cannot make these transformations
 - lack of associativity and distributivity in floating point arithmetic



2.4 Superscalar Processor

- Def: a superscalar processor can issue and execute multiple instructions in one cycle. Instructions are retrieved from a sequential instruction stream and are usually scheduled dynamically
- Benefit: without programming effort, superscalar processor can take advantage of the instruction level parallelism that most programs have
- most modern CPUs are superscalar

2.4.1 Pipelined Functional Units

```
long mult_eg(long a, long b, long c) {
    long p1 = a*b;
    long p2 = a*c;
    long p3 = p1*p2;
    return p3;
}
```

- divide computations into stages
- pass partial computations from stage to stage

- stage i can start on new computation once values passed to $i+1$
- e.g. complete 3 multiplications in 7 cycles, even though each requires 3 cycles

Time							
	1	2	3	4	5	6	7
Stage 1	$a*b$	$a*c$			$p1*p2$		
Stage 2		$a*b$	$a*c$			$p1*p2$	
Stage 3			$a*b$	$a*c$			$p1*p2$