

Lecture 13

<2016-05-25 Wed>

Contents

1	Operating System	1
1.1	Virtual Memory, Page	1
1.1.1	Solution: Virtual Addressing	2
1.1.2	Address Space	3
1.1.3	Why Virtual Memory	4
1.1.4	Address Translation: Page Table	4
1.2	Exception	5
1.2.1	Altering Control Flow	5
1.2.2	Exceptional Control Flow	6
1.2.3	Exceptions	6
1.2.4	Interrupt Vectors	7
1.2.5	Asynchronous/Synchronous Exception	8
1.3	Linking	9
1.3.1	Static Linking	9
1.3.2	Linker	9
1.3.3	Three Kinds of Object Files	10
1.3.4	Static Libraries (.a archives) vs. Shared Libraries (.so)	10
1.3.5	Dynamic Linking	12

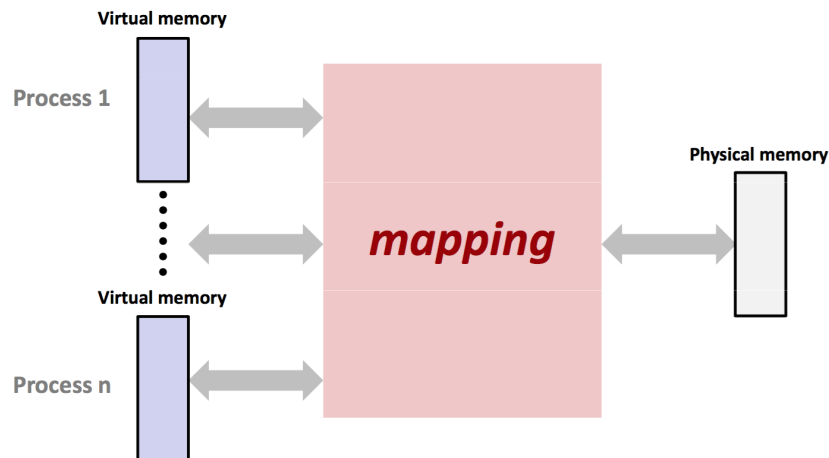
1 Operating System

1.1 Virtual Memory, Page

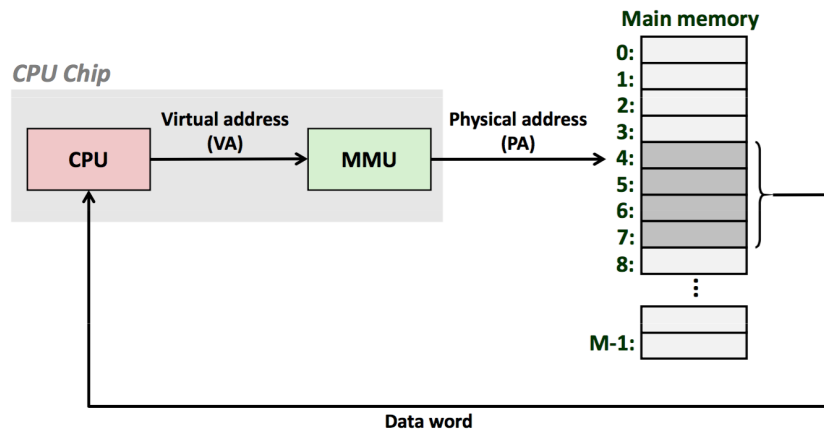
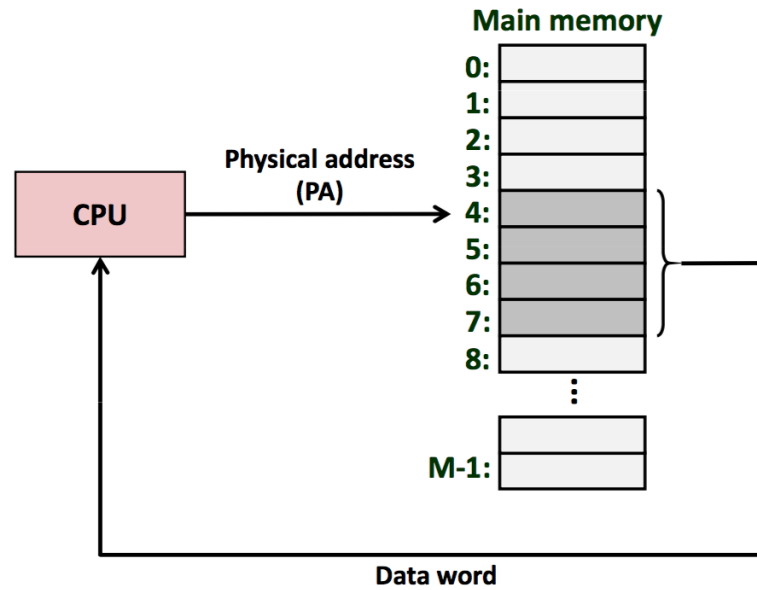
- "programs" refer to virtual memory address
 - `movl (%ecx), %eax`
 - conceptually very large array of bytes

- each bytes have its own address
- actually implemented with hierarchy of different memories
- system provides address space private to particular "program"
- allocation: compiler and run-time system
 - where different program objects should be stored
 - add allocation within single virtual address space
- why virtual memory
 - address does not fit
 - * 64-bit address space(16 exabyte) » physical main memory(few gigabytes)
 - memory management
 - how to protect
 - * how to prevent different processes from using the same physical memory
 - how to share

1.1.1 Solution: Virtual Addressing



- each process gets its own private memory space
- solves the previous problems



1.1.2 Address Space

- virtual address space
 - set of 2^n virtual addresses
- physical address space
 - set of 2^m physical addresses

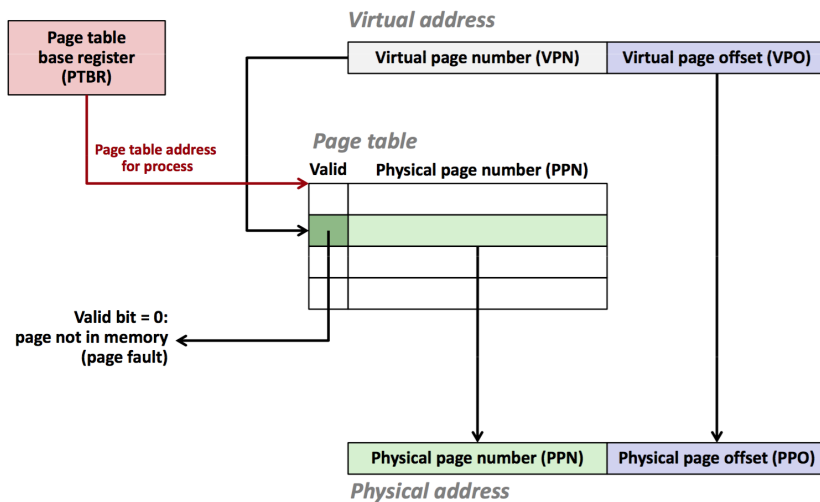
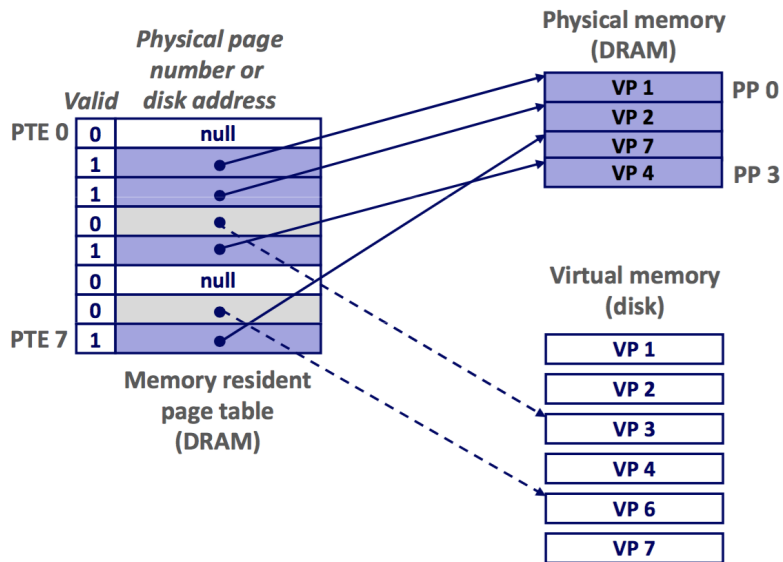
- clean distinction between data (bytes) and their attributes (addresses)
- each object can now have multiple addresses
- every byte in main memory
 - 1 physical address
 - 1 (or more) virtual addresses

1.1.3 Why Virtual Memory

- efficient use of limited main memory (RAM)
 - use RAM as cache for the parts of a virtual address space
 - * some non-cache parts stored on disk
 - * some (unallocated) non-cache parts stored nowhere
 - keep only active of areas of virtual address space in memory
 - * transfer data back and forth as needed
- simplifies memory management for programmers
 - each program gets the same full, private linear address space
- isolates address space
 - one process can't interfere with another's memory
 - * because they operate in different address space
 - user process can't access privileged information
 - * different sections of address spaces have different permissions

1.1.4 Address Translation: Page Table

- a page table is an array of page table entries (PTEs) that maps virtual pages to physical pages



1.2 Exception

1.2.1 Altering Control Flow

- control flow
 - a CPU simply reads and executes a sequence of instructions
 - this sequence is the system's physical control flow

- alter the control flow: 2 mechanisms
 - jumps and branches
 - call and return using the stack discipline
 - both react to changes in program state
 - insufficient for a useful system
 - * difficult for the CPU to react to changes in system state
 - data arrive from a disk or network adapter
 - instructions divides by 0
 - user hits control-c at the keyboard
 - system timer expires
- needs exceptional control flow

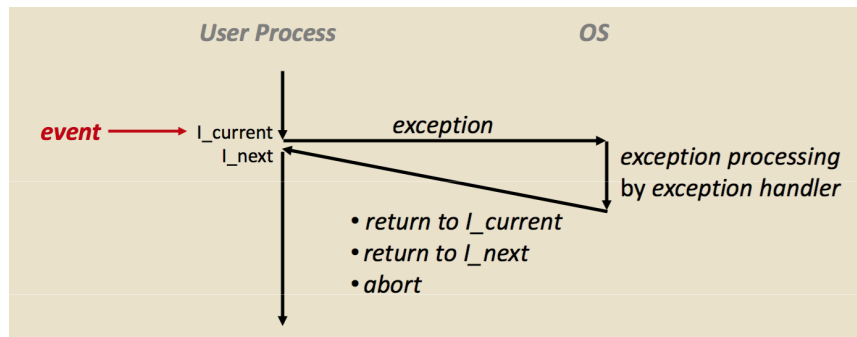
1.2.2 Exceptional Control Flow

- mechanisms for exceptional control flow exists at all level of a computer system
- low level mechanism
 - exceptions
 - * change in control flow in response to system event
 - * combination of hardware and OS software
- high level mechanisms
 - process context switch
 - signals
 - nonlocal jumps
 - implemented by either
 - * OS software: context switch and signals
 - * C language runtime library: nonlocal jump

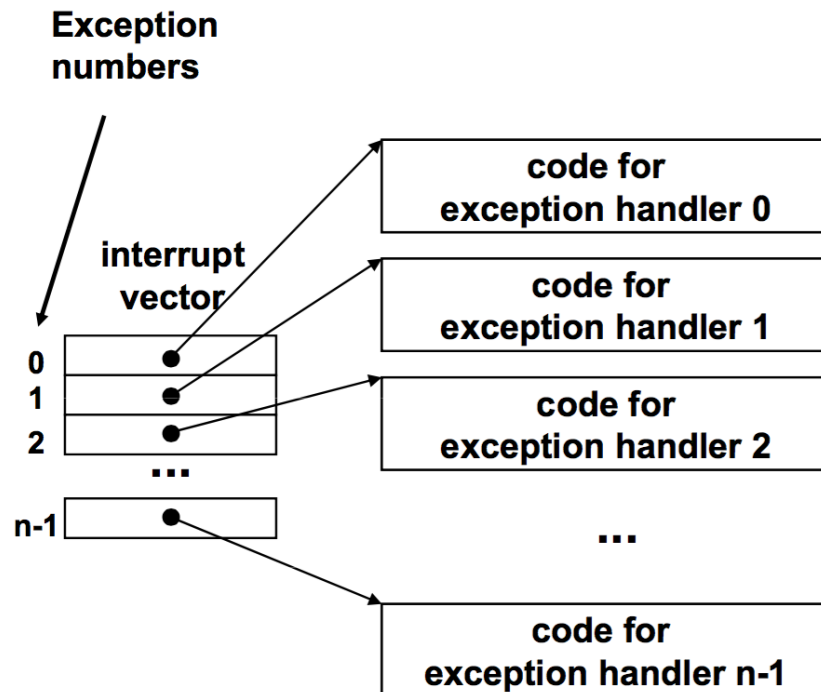
1.2.3 Exceptions

- an exception is a transfer of control to the OS in response to some event
- examples

- divide by 0
- arithmetic overflow
- page fault
- I/O request completes
- Ctrl-C



1.2.4 Interrupt Vectors



- each type of event has a unique exception number k
 - k as an index into jump table (interrupt vector)
- jump table entry k points to an exception handler
 - handler k is called each time exception k occurs

1.2.5 Asynchronous/Synchronous Exception

- asynchronous
 - caused by events external to the processor
 - * indicated by setting the processor's interrupt pin
 - * handler returns to next instruction
 - examples
 - * I/O interrupts
 - hitting Ctrl-C at the keyboard
 - arrival of a packet from a network
 - arrival of a data sector from a disk
 - * hard reset interrupt
 - hitting the reset button
 - * soft rest interrupt
 - hitting Ctrl-Alt-Delete on a PC
- synchronous
 - caused by events that occur as a result of executing an instruction
 - * traps
 - intentional
 - example: system calls, breakpoint traps, special instructions
 - returns control to next instruction
 - * faults
 - unintentional but possibly recoverable
 - example: page fault (recoverable), protection faults (unrecoverable)
 - either re-executes faulting ("current") instruction or aborts
 - * aborts

- unintentional and unrecoverable
- example: parity error, machine check
- aborts current program

1. Read A Disk Sector (asynchronous)

2. Trap Example (synchronous)

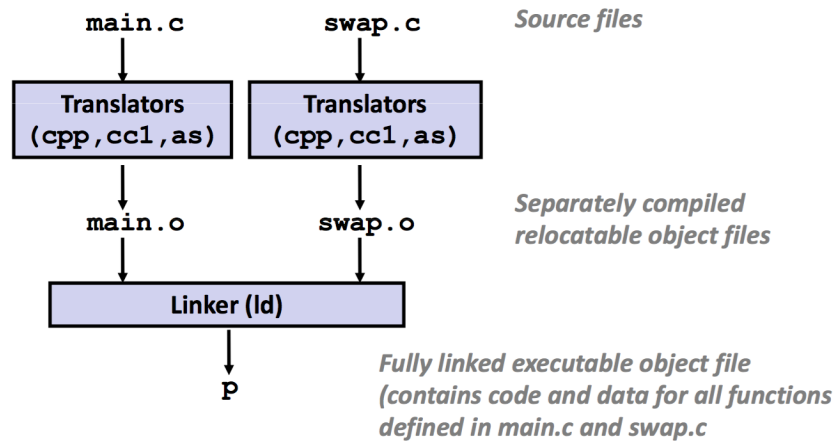
- open a file
 - user call `open` (system call)
 - OS must find or create file, get it ready for reading or writing
 - returns integer file descriptor

3. Fault Example (synchronous)

- memory reference

1.3 Linking

1.3.1 Static Linking



- programs are translated and linked using a compiler driver

1.3.2 Linker

- modularity

- program can be written as a collection of smaller source files, rather than one monolithic mass
- can build libraries of common functions
 - * e.g., math library, standard C library
- efficiency
 - time: separate compilation
 - * change one source file, compile, and then relink
 - * no need to recompile other source files
 - space: libraries
 - * common functions can be aggregated into a single file
 - * yet executable files and running memory images contain only code for the functions they actually use

1.3.3 Three Kinds of Object Files

- relocatable object file (`.o`)
 - contain code and data in a form that can be combined with other relocatable object files to form executable object files
 - each `.o` file is produced from exactly one source file
- executable object file
 - contains code and data in a form that can be copied directly into memory and then executed
- shared object file (`.so`)
 - special type of relocatable object file that can be loaded into memory and linked dynamically, at either load-time or run-time
 - called dynamic link libraries (dll) by windows

1.3.4 Static Libraries (`.a` archives) vs. Shared Libraries (`.so`)

1. Static Libraries

- concatenate related relocated object files into a single file with an index

- enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives
- if an archive member file resolves reference, link into executable
- disadvantage:
 - duplication in the stored executables (every function need stdlibc)
 - duplication in the running executables
 - minor bug fixes of system require each application to explicitly relink

2. Shared Libraries

- object files that contain code and data that are loaded and linked into an application dynamically, at either load-time or run-time
- also called: dynamic link libraries, DLLs, .so files

3. Shared Libraries and Dynamic Linking

- dynamic linking occur when executable is first loaded and run (**load-time linking**)
 - common case for Linux, handled automatically by the dynamic linker
 - standard C library (libc.so) usually dynamically linked
- dynamic linking can also occur after program has begun (**run-time linking**)
 - in unix, this is done by calls to the **dlopen** interface

1.3.5 Dynamic Linking

